

`Rc<T>` and `RefCell<T>`

CAS CS 392: Rust, in Theory and in Practice

Lecture 11

Outline

Reference Counting

Interior Mutability

A Couple Last Things

High Level

Ownership is nice, borrowing can make it ownership nice to work with, but we *still* may want data to have multiple owners:

```
// THIS DOES NOT COMPILE
use crate::List::{Cons, Nil};
enum List {Cons(i32, Box<List>), Nil}
fn main() {
    let a = Cons(5, Box::new(Cons(10, Box::new(Nil))));
    let b = Cons(3, Box::new(a));
    let c = Cons(4, Box::new(a));
}
```

Recall: Reference Counting

Languages like Swift allow for this using *reference counting*:

```
class Stuff {  
    init() {print("allocating")}  
    deinit {print("deallocating")}  
}  
  
var r1 : Stuff? = Stuff()  
var r2 : Stuff? = r1  
var r3 : Stuff? = r2  
  
r1 = nil           // prints:  
r2 = nil           // allocating  
r3 = nil           // deallocating
```

- ▶ The number of references is maintained
- ▶ Once the count goes to zero, the associated memory can be freed

Recall: Lists

```
// THIS DOES NOT COMPILE
use crate::List::{Cons, Nil};
enum List {Cons(i32, Box<List>), Nil}
fn main() {
    let a = Cons(5, Box::new(Cons(10, Box::new(Nil))));
    let b = Cons(3, Box::new(a));
    let c = Cons(4, Box::new(a));
}
```

(Last time) We can build recursive list data types using boxes

Boxes still don't allow for complex reference structures, like two lists with the same tail because **a box must still have a single owner**

The Reference Counting Type

```
use crate::List::{Cons, Nil};
use std::rc::Rc;
enum List {Cons(i32, Rc<List>), Nil}
fn main() {
    let a = Rc::new(Cons(5,
                        Rc::new(Cons(10,
                                    Rc::new(Nil))));
    let b = Cons(3, Rc::clone(&a));
    let c = Cons(4, Rc::clone(&a));
}
```

We can get these kinds of structures with the *reference counted smart pointers*

`Rc<T>` is like a box, except that it's possible to clone the reference *without cloning the data itself*

`Rc::clone(&a)` just increments the reference count

Aside: Deref and Reference Counting

`Rc<T>` implements the `Deref` trait:

```
let x : Rc<i32> = Rc::new(5);  
let x_ref : &Rc<i32> = &x;  
let x_ref2 : &i32 = &x;  
let _ : Rc<i32> = Rc::clone(x_ref);  
// let _ : Rc<i32> = Rc::clone(x_ref2);
```

`Rc::clone` has the following signature:

```
fn clone(&self) -> Rc<T, A>
```

The point: How a reference is coerced can be determined by the type annotation

Rc::clone vs. .clone()

We can also just use `.clone()`:

```
let x : Rc<i32> = Rc::new(5);  
let _ = x.clone();
```

The Rust convention is to use `Rc::clone` to distinguish *shallow clones* from *deep clones*

References Counts

Clones increment the reference count:

```
impl<T: ?Sized, A: Allocator + Clone> Clone for Rc<T, A> {  
    fn clone(&self) -> Self {  
        unsafe {  
            self.inner().inc_strong();  
            Self::from_inner_in(self.ptr,  
                               self.alloc.clone())  
        }  
    }  
}
```

We can use `Rc::strong_count` to see what the current count is:

```
fn main() {  
    let x = Rc::new(5); let y = x.clone();  
    assert_eq!(2, Rc::strong_count(&x)); drop(y);  
    assert_eq!(1, Rc::strong_count(&x));  
}
```

Immutability

Reference counted smart pointers are immutable:

```
// THIS DOES NOT COMPILE  
let x : Rc<i32> = Rc::new(5);  
*x = 3;
```

If we have multiple references, we don't want to be able to mutate them!

Reference Counting and Concurrency

We also have `Arc<T>` for (multi-)thread safe ("atomic") reference counting:

```
let counter = Arc::clone(&counter);  
let handle = thread::spawn(move || {  
    let mut num = counter.lock().unwrap();  
    *num += 1;  
});
```

`Rc<T>` should only be used the in single-threaded settings (generally better performance)

Outline

Reference Counting

Interior Mutability

A Couple Last Things

RefCell<T>

```
fn main() {  
    let x = RefCell::new(5);  
    let mut y : RefMut<'_, i32> = x.borrow_mut();  
    *y += 1; drop(y);  
    println!("{}", x.borrow())  
}
```

RefCell<T> owns the data it holds like a Box<T>

We get a compile time error if we create multiple mutable references to a Box<T>

We get a **runtime error** if we create multiple mutable references to a RefCell<T>

The Takeaway: We can mutate a value in a RefCell<T> even when it is immutable

Trade-offs

Compile-time errors give us better assurance that our code is correct

There are memory-safe operations that are not allowed by the borrow checker

The Rust compiler is conservative. *Annoyance is better than catastrophe*

Comparisons

| smart pointer | # owners | mut ref allowed? | mut ref checked? |
|---------------|----------|------------------|------------------|
| Box<T> | one | yes | compile-time |
| Rc<T> | many | no | N/A |
| RefCell<T> | one | yes | run-time |

Example: Mock Objects

There's a nice example in the text on this:

```
struct MockMessenger {sent_messages: RefCell<Vec<String>>},  
impl MockMessenger {  
    fn new() -> MockMessenger {  
        MockMessenger {sent_messages: RefCell::new(vec![]),}  
    }  
}  
impl Messenger for MockMessenger {  
    fn send(&self, message: &str) {  
        self.sent_messages  
            .borrow_mut().push(String::from(message));  
    }  
}
```

The rough idea: If we *have* to work with something by (immutable) reference (e.g., because of a trait), we can have mutable state in a RefCell

Reference Counting + Interior Mutability

We can combine reference counting and interior mutability to have mutable values in things like lists:

```
enum List {Cons(Rc<RefCell<i32>>, Rc<List>), Nil}
fn main() {
    let value = Rc::new(RefCell::new(5));
    let a = Rc::new(Cons(Rc::clone(&value), Rc::new(Nil)));
    let b = Cons(Rc::new(RefCell::new(3)), Rc::clone(&a));
    let c = Cons(Rc::new(RefCell::new(4)), Rc::clone(&a));
    *value.borrow_mut() += 10;
}
```

Outline

Reference Counting

Interior Mutability

A Couple Last Things

Reference Cycles

Combining `RefCell<T>` and `Rc<T>` can create *reference cycles*:

```
let a = Rc::new(Cons(5, RefCell::new(Rc::new(Nil))));  
let b = Rc::new(Cons(10, RefCell::new(Rc::clone(&a))));  
if let Some(link) = a.tail() {  
    *link.borrow_mut() = Rc::clone(&b);  
}
```

Reference cycles can lead to leaked memory (it's impossible to bring the reference count down to 0)

This means /memory leaks are "safe" in Rust

An Implementation of Rc<T>

```
pub struct Rc<
    T: ?Sized,
    A: Allocator = Global,
> {
    ptr: NonNull<RcInner<T>>,
    phantom: PhantomData<RcInner<T>>,
    alloc: A,
}
// ...
fn clone(&self) -> Self {
    unsafe {
        self.inner().inc_strong();
        Self::from_inner_in(self.ptr, self.alloc.clone())
    }
}
```

What the heck is the NonNull? And what is PhantomData? And is it okay to use this unsafe code?